Defensive metrics: Cressey

What the defender gets, and what is and is not measured. Numbers come from the evaluation runs in the repo.

Time saved for the analyst

Not measured. The repo has no baseline of analyst time. What is measured is machine time: time to stop a call, p50 629 ms with live Jev and the calls back to back (236 ms at the stage's pace, one run, 10-10); three fix guides in 9 to 14 s (10-10 demo path).

Estimate (assumptions, not data):

Step per held incident	Without Cressey	With Cressey	Saved
Find what the agent did and why	15 to 30 min	2 to 5 min	10 to 25 min
Find where the misused secret came from	10 to 30 min	under 1 min	10 to 30 min
Write a fix guide per owner	30 to 60 min	about 5 min to review	25 to 55 min
Decide the hold	5 min	2 min	about 3 min
Total	**about 1 to 2 h**	**about 10 to 15 min**	**about 50 to 100 min**

The "without" column is a guess at a careful manual triage. Replace both columns with a timed test on the synthetic incidents. Not counted: false-alarm review time on real traffic, and approver time on writes held without a security reason.

Precision of detections and conclusions

Check	Result	Limit
Holds on normal sessions	**0/6** held by mistake	6 synthetic sessions, written by the team
Incident sessions held	**8/8**, each at the call it names	the team wrote the incidents to match the rules
Write judgment at 0.5	30/30 on labelled writes (15 yes, 15 no)	labels are ours
Report check at 0.8	**5/5** precision, 5/10 recall	20 reports; misses listed in evaluation
Sourced guide steps	100% citation validity	checks that a step cites a retrieved passage, not that the advice is good
Retrieval hit@3 after re-ranking	30 to 31 of 31	33 team-written questions

These are weak evidence on a set this small. They show the rules and triangle work on our own sessions, not a false-positive rate for real traffic. The held-out set from outside the team is not written.

Coverage

Framework	Covered	Coverage in the repo
MITRE ATT&CK	T1485, T1552.001, T1078	Mapped in the pitch. The destructive-SQL and trap-key sessions exercise T1485 and T1552.001
MITRE ATLAS v5.6.0	8 techniques: AML.T0053, T0101, T0055, T0083, T0051, T0080, T0110, T0086	IDs checked against the ATLAS reference file. The guide corpus holds 35 ATLAS mitigations
OWASP Agentic	ASI01, ASI02, ASI03	**The OWASP corpus file is not in the repo** (download by hand), so no OWASP guide is checked against it

Gaps: ATLAS and ATT&CK are hand mappings, and nothing tests them beyond the sessions named. Coverage holds only for traffic through Cressey.

Summary

Measured: precision on the team's synthetic set (0 false holds, 8/8 incidents held) and machine speed. Not measured: analyst time saved. Coverage is a hand mapping, and the OWASP corpus is not loaded.